

Linear Regression

August 27, 2020

1 Linear Regression

Written by: Paul Treve (Email: trevepaul2@gmail.com)

Note: Based on the data provided, an accurate hypothesis should be $h_{\theta}(x^{(i)}) = X\theta = 1.5 + 2x_1^{(i)} + x_2^{(i)}$

$$\theta = \begin{pmatrix} 1.5 \\ 2 \\ 1 \end{pmatrix}$$

```
[1]: import numpy as np
import pandas as pd
from matplotlib import pyplot as plt

%matplotlib inline
```

1.1 Finding the parameters of hypothesis using gradient descent

1.1.1 Handling the data from the csv file into arrays I can work with

```
[2]: # loading the data from the csv file
raw_data = pd.read_csv('data.csv')

m = len(raw_data) # number of training examples
n = 2 # number of features

# assuming the hypothesis is of the form X*theta where X is an m by n matrix
# and theta, an n by 1 matrix
feature1 = np.array(raw_data['X']) #x1
feature2 = np.array(raw_data['Y']) #x2

X = np.array([
    np.ones(m), # x0 is 1
    feature1,
    feature2
]).transpose()

Y = np.array([[z] for z in np.array(raw_data['Z'])]) # correct results
```

```
# initializing theta as an n by 1 matrix made up of only ones
theta = np.array([[1] for _ in range(3)])
```

What is X? X is the design matrix of m rows and $(n + 1)$ columns. An $m \times (n + 1)$ matrix with each row representing a training example of the dataset.

```
[3]: X[:5] # first 5 rows of X
```

```
[3]: array([[1., 2., 3.],
           [1., 1., 2.],
           [1., 4., 1.],
           [1., 3., 2.],
           [1., 2., 1.]])
```

What is Y? Y is an m dimensional vector. An $m \times 1$ matrix with each element corresponding to the expected result of a specific training example of the dataset.

```
[4]: Y[:5] # first 5 elements of Y
```

```
[4]: array([[ 8.5],
           [ 5.5],
           [10.5],
           [ 9.5],
           [ 6.5]])
```

1.1.2 The Hypothesis - $h_{\theta}(x^{(i)})$

Instead of taking each feature, $x_n^{(i)}$, one at a time in a training example to multiply with its corresponding θ_j , a more conservative way to do this is to vectorize the θ . Making θ an $(n + 1)$ dimensional vector, with the first element being the constant term in the expression $h_{\theta}(x^{(i)}) = \theta_0 + \theta_1 x_1^{(i)} + \theta_2 x_2^{(i)} + \dots + \theta_n x_n^{(i)}$ and the rest of the elements being the coefficients of $x_1^{(i)}, x_2^{(i)}, \dots, x_n^{(i)}$ where n is the number of features of each training example. Multiplying the design matrix, X (an $m \times (n + 1)$ matrix), by θ will produce an m dimensional vector (an $m \times 1$ matrix). This vector contains the predictions of the hypothesis.

$$h_{\theta}(x^{(i)}) = X\theta$$

```
[5]: def hypothesis():
      return X @ theta # matrix product of X and theta
```

After some trials, a learning rate of 0.01 is enough for the gradient descent.

```
[6]: # the learning rate
alpha = 0.01
```

1.1.3 The Cost Function for Linear Regression

$$J(\theta) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2$$

```
[7]: def cost():  
      return (1/(2*m))*(((hypothesis() - Y) ** 2).sum())
```

1.1.4 The Gradient Descent

$$\theta_j := \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta)$$

But since I have decided to vectorize the θ , I can think of the gradient descent as:

$$\theta := \theta - \alpha \delta$$

where δ is representing the partial derivative of each of elements of θ . δ will be an $(n+1)$ dimensional vector given by:

$$\delta = \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x^{(i)}$$

The function in the cell below does the update of the theta as it should be done - simultaneously.

```
[8]: def new_theta():  
      x0 = np.array([u for u in X.transpose()[0]])  
      x1 = np.array([u for u in X.transpose()[1]])  
      x2 = np.array([u for u in X.transpose()[2]])  
      # computing the elements of the delta vector  
      delta0 = (1/m)*(((hypothesis() - Y)*x0).sum())  
      delta1 = (1/m)*(((hypothesis() - Y)*x1).sum())  
      delta2 = (1/m)*(((hypothesis() - Y)*x2).sum())  
      # computing the new values for the elements of theta  
      return np.array([  
          [theta[0, 0] - (alpha*delta0)],  
          [theta[1, 0] - (alpha*delta1)],  
          [theta[2, 0] - (alpha*delta2)]  
      ])
```

A while loop to update the theta and keep the gradient descent going until the cost is minimal

```
[9]: costs = np.array([])  
      number_of_steps = 0  
      while True:  
          # keeping track of the number of steps and the gradient descent at each step  
          # this is to help plot a graph of cost against number of steps  
          # to show that the gradient descent actually worked  
          number_of_steps += 1  
          costs = np.append(costs, cost())
```

```

# calculation of accuracy of model in terms of MSE
mean_sq_error = 2*cost()
# making sure the cost is minimum
if cost() < 10**(-20):
    print("Success.")
    break
# updating the theta simultaneously
theta = new_theta()

```

Success.

```

[10]: # printing out results on the screen
print("Results\n-----")
print(f"Theta:\n {theta}\nMean Squared Error(MSE): {mean_sq_error: .15f}")

```

Results

Theta:

[[1.5]

[2.]

[1.]]

Mean Squared Error(MSE): 0.0000000000000000

```

[11]: # This is a plot of costs against number of steps taken in gradient descent.
# The nature of the graph determines the success of the gradient descent.
plt.subplot(221)
plt.plot(range(number_of_steps), costs, '.-r')
plt.xlabel('Number of steps')
plt.ylabel('Costs')
plt.title('A plot of Costs against the number of steps')
plt.subplot(222)
plt.plot(range(50), costs[:50], '.-g')
plt.xlabel('Number of steps')
plt.ylabel('Costs')

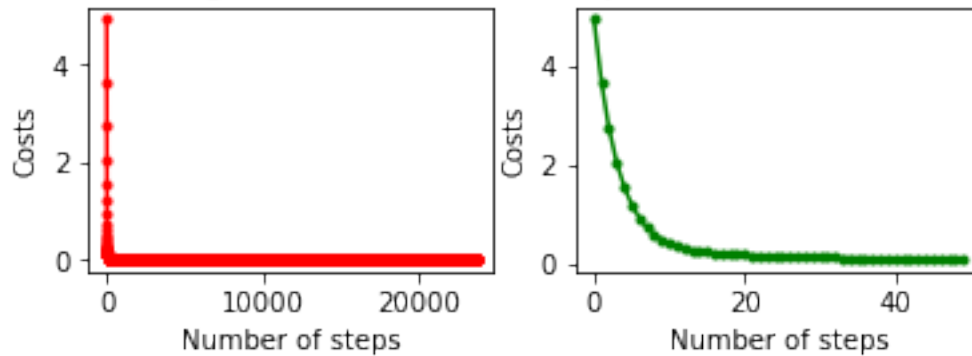
```

```

[11]: Text(0, 0.5, 'Costs')

```

A plot of Costs against the number of steps



1.2 Finding the parameters of hypothesis using the Normal Equation

1.2.1 Handling the data from the csv file into arrays I can work with

```
[12]: data = pd.read_csv('data.csv')
x_feature = np.array(data['X'])
y_feature = np.array(data['Y'])

Y = np.array([[z] for z in data['Z']])

X = np.array([
    np.ones(len(x_feature)),
    x_feature,
    y_feature
]).transpose()
```

What is X? X is the design matrix of m rows and $(n + 1)$ columns. An $m \times (n + 1)$ matrix with each row representing a training example of the dataset.

```
[13]: X[:3] # first 3 rows of X
```

```
[13]: array([[1., 2., 3.],
            [1., 1., 2.],
            [1., 4., 1.]])
```

What is Y? Y is an m dimensional vector. An $m \times 1$ matrix with each element corresponding to the expected result of a specific training example of the dataset.

```
[14]: Y[:3] # first 3 elements of Y
```

```
[14]: array([[ 8.5],
            [ 5.5],
```

```
[10.5]])
```

1.2.2 The normal equation

The normal equation is given by:

$$\theta = (X^T X)^{-1} X^T Y$$

```
[15]: theta = np.linalg.inv(X.transpose() @ X) @ X.transpose() @ Y
```

```
[16]: # printing out results on the screen
print("Results\n-----")
print(f"Theta:\n {theta}")
```

Results

Theta:
[[1.5]
[2.]
[1.]]

1.3 The Data used for the above codes

```
[17]: Data = pd.read_csv('data.csv')
```

```
[18]: Data.shape
```

```
[18]: (15, 3)
```

```
[19]: Data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 15 entries, 0 to 14
Data columns (total 3 columns):
#   Column  Non-Null Count  Dtype
---  -
0    X      15 non-null        int64
1    Y      15 non-null        int64
2    Z      15 non-null        float64
dtypes: float64(1), int64(2)
memory usage: 488.0 bytes
```

```
[20]: Data.describe()
```

```
[20]:
```

	X	Y	Z
count	15.000000	15.000000	15.000000
mean	2.466667	2.533333	8.966667
std	1.060099	1.187234	2.503331

min	1.000000	1.000000	4.500000
25%	2.000000	1.500000	7.500000
50%	2.000000	3.000000	8.500000
75%	3.000000	3.500000	10.500000
max	4.000000	4.000000	13.500000

1.3.1 The full data

[21]: Data

```
[21]:
```

	X	Y	Z
0	2	3	8.5
1	1	2	5.5
2	4	1	10.5
3	3	2	9.5
4	2	1	6.5
5	2	2	7.5
6	1	4	7.5
7	3	3	10.5
8	1	1	4.5
9	3	4	11.5
10	4	4	13.5
11	3	1	8.5
12	2	3	8.5
13	2	4	9.5
14	4	3	12.5